



Expand All

Collapse All

## Activation function

### ▼ Which Math library functions may I use?

Use `Math.exp()` and `Math.abs()`. You may also use the constants `Math.POSITIVE_INFINITY`, `Math.NEGATIVE_INFINITY`, and `Math.NaN`.

### ▼ How do I detect whether $x$ is NaN?

Use `Double.isNaN(x)`. Do not write code like `(x == Double.NaN)`, as that expression evaluates to `false` for all values of  $x$ , including NaN.

### ▼ My `softsign()` function return NaN when $x$ is positive infinity (or negative infinity) instead of 1 (or -1). Why is this?

You are dividing infinity by infinity, which leads to NaN (not a number). Include special cases for when  $x$  equals either `Double.POSITIVE_INFINITY` or `Double.NEGATIVE_INFINITY`,

### ▼ My `tanh()` functions return NaN when $x$ is larger than 750 (or smaller than -750) instead of 1 (or -1). Why is this?

The term  $e^x$  will be larger than the largest floating-point number. As a result, you end up dividing infinity by infinity, which leads to NaN. You may use the fact that if  $x \geq 20$ , then  $\tanh(x)$  should return 1.0; if  $x \leq -20$ , then  $\tanh(x)$  should return -1.0.

## Divisors

### ▼ May I use recursion to implement `gcd()`?

That's a fine idea, but please use a loop this week. We'll learn about *recursion* next week.

### ▼ How do I compute `gcd()` and `lcm()` when one (or both) of the arguments are negative?

It follows from the definitions that  $\gcd(a, b) = \gcd(|a|, |b|)$  and  $\text{lcm}(a, b) = \text{lcm}(|a|, |b|)$ .

### ▼ Why does Euclid's algorithm correctly compute the greatest common divisor of two integers?

It is based on the observation that any common divisor  $d$  of the integers  $a$  and  $b$  must also divide the differences  $a - b$ ,  $a - 2b$ ,  $a - 3b$ , ..., which includes  $a \% b$ . Here's a [full explanation](#).

### ▼ How many iterations does Euclid's algorithm take to compute the greatest common divisor of $a$ and $b$ ?

As you would expect, the number of iterations depends on  $a$  and  $b$ . It turns out that the smallest pair of integers that causes Euclid's algorithm to take  $n$  iterations are the consecutive Fibonacci numbers  $F_{n+2}$  and  $F_{n+1}$ . This implies that the number of steps is at most  $5 \log_{10} b$ , i.e., at most 5 times the number of decimal digits in the smaller number.

### ▼ Why do we use the conventions that $\gcd(0, 0) = 0$ and $\text{lcm}(0, k) = 0$ ?

These are consistent with the more technical definitions of  $\gcd$  and  $\text{lcm}$  as the *meet* (greatest lower bound) and *join* (least upper bound) in the [lattice](#) of divisibility.

### ▼ Are there faster algorithms for computing Euler's totient function?

Yes. Computing the totient function of an integer  $n$  is computationally equivalent to factoring  $n$ . Specifically, if the distinct prime factors of  $n$  are  $p_1, p_2, \dots, p_r$ , then  $\phi(n) = n \left(1 - \frac{1}{p_1}\right) \left(1 - \frac{1}{p_2}\right) \dots \left(1 - \frac{1}{p_r}\right)$ . So, while computing the totient function is believed to be a computationally intractable problem for large  $n$ , it can be done substantially faster than brute-force approach adopted in this assignment.

## Audio collage

### ▼ What should each method do if one (or more) of the input samples is not between -1 and +1?

That's ok. Just handle it as usual. While you should not play samples whose absolute value is greater than 1, it's fine to manipulate such values along the way.

### ▼ Can `amplify()` or `mix()` produce samples whose absolute value is larger than 1, even if all of the input samples are between -1 and +1?

Yes. While you should not play samples whose absolute value is greater than 1, it's fine to produce them as intermediate results.

### ▼ How can I convert an audio file into an appropriate format for use with `StdAudio`?

Convert it to a WAV file. Be sure to use 16-bit audio, monaural, and a sampling rate of 44,100 Hz. You may use an online conversion utility, such as [Online-Convert](#).

### ▼ What is a WAV file?

It is popular file format for storing raw and uncompressed audio data.

### ▼ The `changeSpeed()` function changes not only the speed of the sound, but also the pitch? Is there a way to change the speed without affecting the pitch?

Speeding up the sound using resampling raises the pitch and leads to the [Chipmunk effect](#). More sophisticated [time stretching](#) techniques are preferred in practice (such as when watching Coursera videos at 1.5× or 2× speed) because they change the speed but not the pitch.

### ▼ Which other audio effects might I want to implement?

- *clamp*: round all samples greater than +1 to +1; round all samples less than -1 to -1.

- *normalize*: rescale a sound so that all values are between  $-1$  and  $+1$ .
- *cut*: extract a contiguous subarray from a given sound.
- *trim*: remove leading / trailing sequence of samples that are 0 (or nearly 0).
- *loop*: repeat a given sound a specified number of times.
- *mirror*: concatenate a sound with its reverse.
- *hip-hop*: increase speed of a sound; mirror it; then loop it.
- *echo, delay, reverb*: add a time-delayed version of a sound to itself, attenuated by a given factor.
- *fade-in, fade-out*: gradually increase/decrease the volume at the beginning/end of a sound.
- *crossfade*: fade-out first sound; fade-in second sound; overlap.
- *tremolo*: create a trembling effect by modulating the amplitude up and down.

You could also synthesize your own sounds by creating a *sine wave*, *square wave*, *triangle wave*, or *sawtooth wave* of a given amplitude, frequency, and duration.